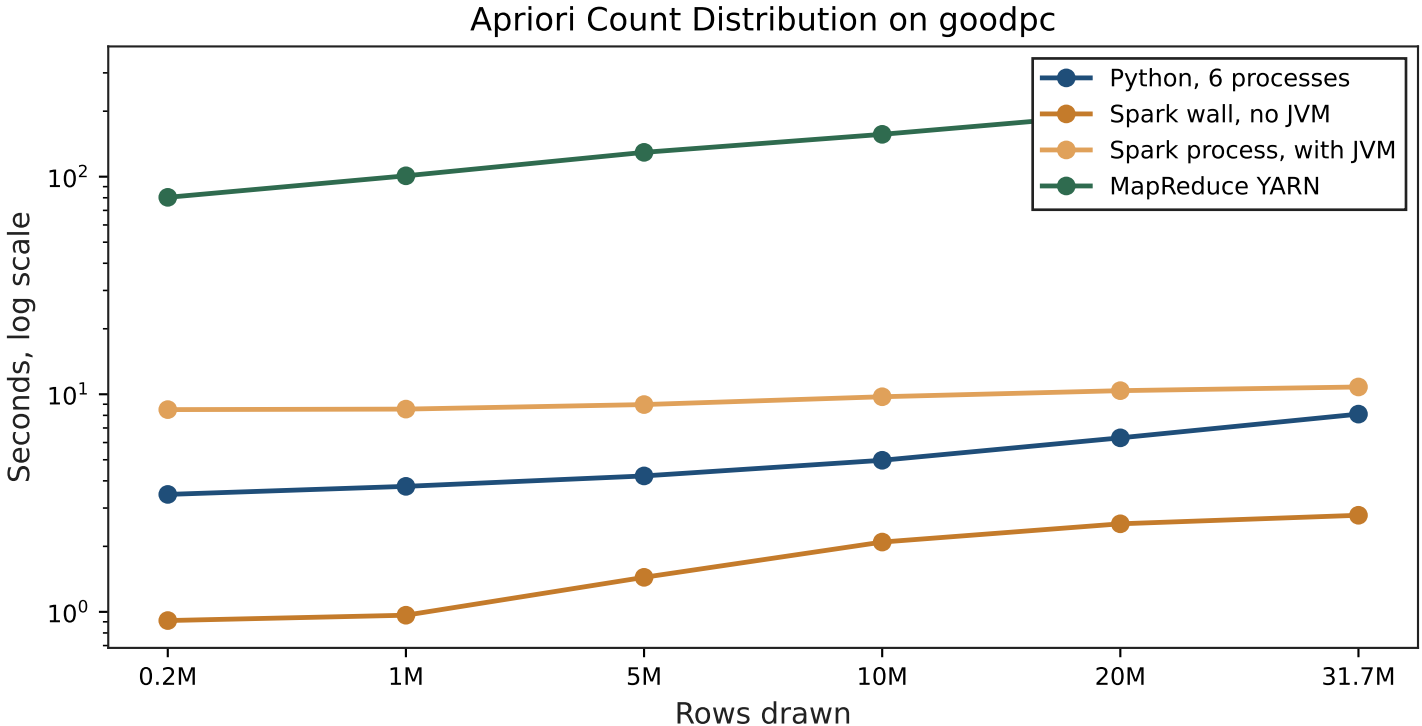

goodpc, Ryzen 5 5600X, 31.89 GB. Support 0.05, confidence 0.45, lift 1.15, itemsets up to 3.

Spark is local[*]. Spark wall is the count time. Spark process adds JVM startup.

MapReduce wall is the three count jobs and excludes HDFS put.

Spark wall is faster than Python at every measured size. Spark process is slower at every measured size.